

# **University of Missouri-Kansas City**

## **School of Computing and Engineering**



**ICP:** ICP-04

**Course Name:** Big Data Analytics and Applications  
**Course ID:** COMP-SCI 5542

**Semester /Session :** Spring 2022

**Student ID : Student Name**

16334245 Mustavi Islam

16321217 Keenan Flynn

### **Description of the Problem :**

ICP 4 focused on the tuning of neural network hyperparameters. Hyperparamaters include the toggleable features of each layer in the network, number of neurons, activation function, optimizer, batch size, epochs and others. Hyperparameter tuning can cause subtle differences in the accuracy of your model and those differences can be both positive and negative. Tuning can change the learning rate, variance, computational cost, and metrics so it is important to understand why this happens.

### **Description of Solution :**

1. The model provided in the source code for the ICP was adapted for the CIFAR-10 dataset.
  - a. The data was already cleaned and split into training and testing sets.
  - b. The labels were transformed using one hot encoding.
  - c. The same Keras sequential model from the source code was then used, but with num\_classes = 10.
2. Tuning the convolutional layers.
  - a. The number of convolutional filters in each Conv2D() layer was increased. Because our images have relatively few dimensions, it is very important that we extract all the features we can. We found through experimentation that we can increase the number of filters for the cost of longer computational time. These filters increase the amount of data that we have about an image and thus increase the accuracy of the model.
  - b. We are currently using 4x the number of filters than what was in the source code. We found this helped bump the model to 80% accuracy
3. Tuning Epochs and Batch size
  - a. Epochs and Batch size were found to be correlated in the performance of the model.
  - b. Initially, a low batch size and low epoch size were used. A low batch size increases learning rate, but also the variance and thus loss of the model. To make sure this loss was not too great, a low number of epochs were used.
  - c. After experimentation, a higher batch size and higher number of epochs were used. Higher batch size decreases the learning rate because the learning of the model occurs after each batch. Increasing batch size decreases the chance for outliers to have a big effect and thus decreases variance and loss. To account for the slow learning rate, a high number of epochs needs to be used.
  - d. We are currently using a batch size of 200 and 35 epochs. We tested all the way up to 50 epochs but found this gave marginal improvement at the cost of marginal deterioration of loss score.

#### 4. Tuning activation functions

- a. The activation function provides non-linearity to the model which separates a neural network problem from a regression problem.
- b. The hidden layers use ‘Relu’ as the activation function. Relu is a good activation function for the hidden layers because it acts like a switch. When a feature is not wanted or needed, Relu ‘turns it off’. When a feature is wanted, Relu is simply the identity function which saves computational cost. For the hidden Dense layer, the activation was changed from softmax to Relu
- c. The output layer is where the model chooses which label to give to a piece of data. This layer was changed to have an activation of ‘softmax’. Softmax outputs a vector of decimals that sum to equal 1. This creates a set of probabilities that the model can choose from. Softmax is a good activation function for multiclass classification problems because of this property.

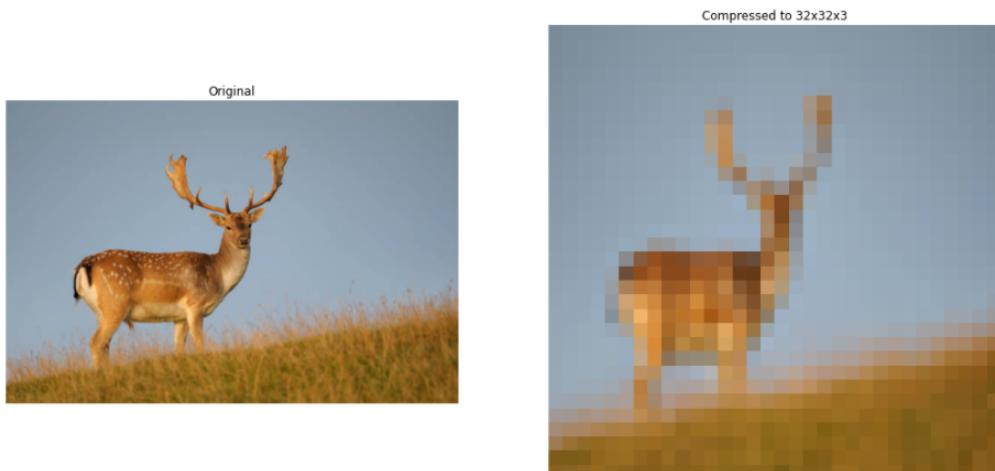
#### 5. The final part of this ICP was to test the model with random images.

- a. These images were scraped from stock photo repos on the web. We created a `loadImage()` function that received a URL as input and output a resized 32x32 image in the form of a numpy array. This function also plots the image so you can visualize the transformation.

#### 6. Results

- a. The model classified the validation data at around 80%. This percentage was reflected by how it classified our random images. It classified 4 out of 5 pictures on a typical run. The random images used were particularly unique images and it is impressive that the model captured these classifications. The images we used are pictured below.

Model correctly classified as deer



Model correctly classified as cat

Compressed to 32x32x3



Model incorrectly classified as truck

Compressed to 32x32x3

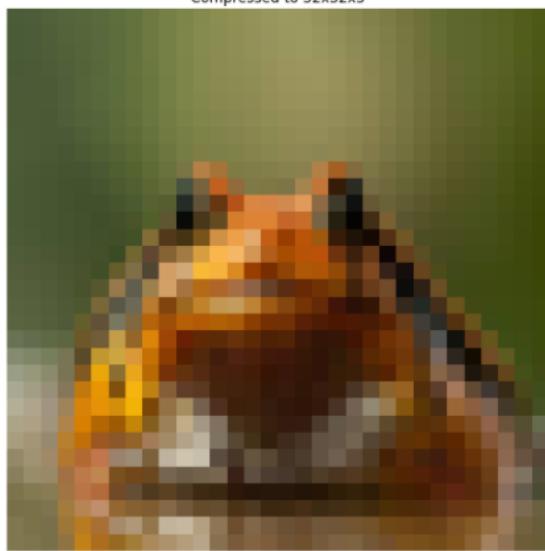


Model correctly classified as ship

Compressed to 32x32x3



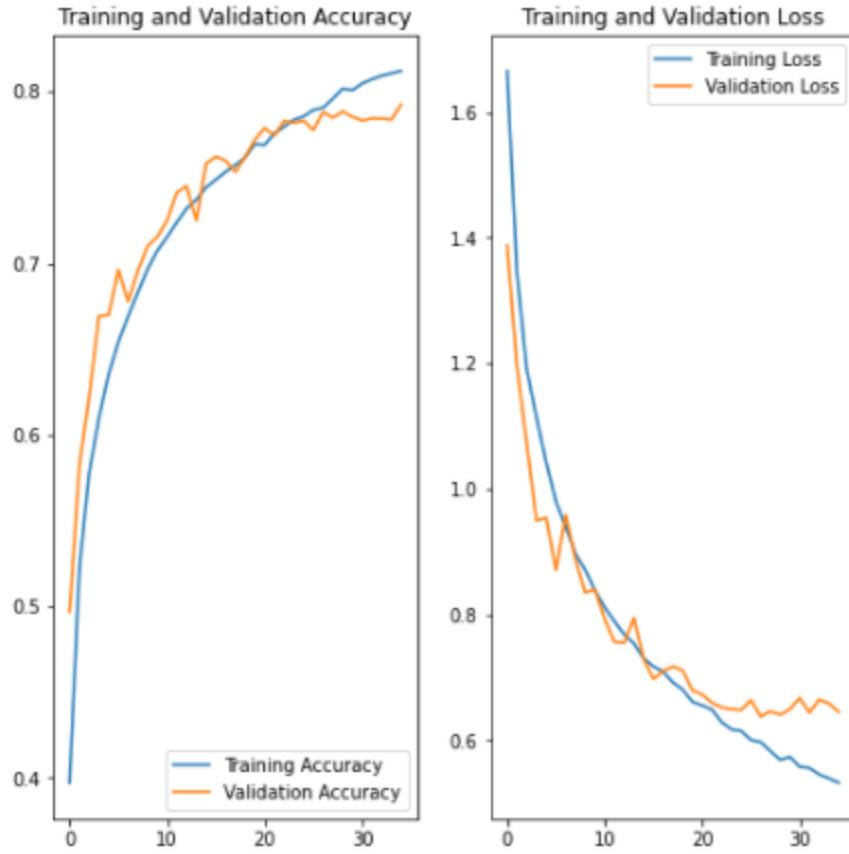
Model correctly classified as frog



**Challenges:**

The biggest challenge this ICP faced was the tuning of hyperparameters. In industry, a grid search can be used to automate this process, but it was informative for us to manually change each hyperparameter and see the effect it had on the model's performance. We tried optimizing many parameters such as filter size, optimizer, and loss function.

We also learned how changing 1 hyperparameter can cause the model to throw errors. One example of this was when we changed the output layer to have an activation function of 'softmax'. Our model started to give a warning, which we had to fix by modifying the loss function. The Categorical cross entropy function had a parameter of 'from\_logits=True'. This was removed because the softmax function does not output logits.



The model classified the validation set with ~80% accuracy. While we could have improved the accuracy further by adding new layers to the model, we decided to solely focus on the tuning of hyperparameters.

### **Learning Outcomes :**

1. We learned how activation functions work.
2. We learned what a grid search was.
3. We learned how to tune a model without adding layers.
4. We learned how to modify hyperparameters to get the best performance.
5. We learned how loss was affected by batch and epoch size.

### **Resources:**

<https://machinelearningmastery.com/how-to-develop-a-cnn-from-scratch-for-cifar-10-photo-classification/> -> We used this site for references on how to validate a model with independent images.

Video Link: <https://youtu.be/dci2wmBiqBs>

## Screenshots :

ICP5-16334245.ipynb

```
[1] import numpy
import tensorflow as tf
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, BatchNormalization, Activation
from keras.layers.convolutional import Conv2D, MaxPooling2D
from keras.constraints import maxnorm
from keras.utils import np_utils
from keras.datasets import cifar10
import PIL
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
```

# Set random seed for purposes of reproducibility  
seed = 21

Now let's load in the dataset. We can do so simply by specifying which variables we want to load the data into, and then using the load\_data() function:

```
[3] # loading in the data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170500996/170498071 [=====] - 3s 0us/step
178508288/170498071 [=====] - 3s 0us/step
```

```
[4] # summarize loaded dataset
print('Train: X=%s, y=%s' % (X_train.shape, y_train.shape))
print('Test: X=%s, y=%s' % (X_test.shape, y_test.shape))

Train: X=(50000, 32, 32, 3), y=(50000, 1)
Test: X=(10000, 32, 32, 3), y=(10000, 1)
```

```
[5] from matplotlib import pyplot as plt
plt.figure(figsize=(8, 8))
# plot first few images
for i in range(9):
    # define subplot
    plt.subplot(3, 3, 1 + i)
```

0s completed at 1:54 AM

ICP5-16334245.ipynb

```
[7] X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
```

[8] `#print(X_train[0])`

The Numpy command to\_categorical() is used to one-hot encode. This is why we imported the np\_utils function from Keras, as it contains to\_categorical().

We also need to specify the number of classes that are in the dataset, so we know how many neurons to compress the final layer down to:

```
[9] # one hot encode outputs
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
class_num = y_test.shape[1]
```

[10] `print(X_train.shape[1:])`  
(32, 32, 3)

[11] `img_height=32  
img_width=32`

```
[12] data_augmentation = keras.Sequential(
        [
            layers.experimental.preprocessing.RandomFlip("horizontal",
                input_shape=(img_height,
                            img_width,
                            3)),
            layers.experimental.preprocessing.RandomRotation(0.1),
            layers.experimental.preprocessing.RandomZoom(0.1),
        ]
    )
```

colab.research.google.com

ICP5-16334245.ipynb - Colaboratory

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[13] num_classes = 10
model = Sequential([
    data_augmentation,
    layers.experimental.preprocessing.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    #layers.Dropout(0.2),
    #layers.BatchNormalization(),
    layers.Conv2D(128, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    #layers.Dropout(0.2),
    #layers.BatchNormalization(),
    layers.Conv2D(256, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Dropout(0.2),
    #layers.BatchNormalization(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes, activation='softmax')
])

[14] model.compile(optimizer='adam',
                    loss=tf.keras.losses.CategoricalCrossentropy(),
                    metrics=['accuracy'])

[15] model.summary()
Model: "sequential_1"
Layer (type)          Output Shape       Param #
sequential (Sequential)    (None, 32, 32, 3)           0
rescaling (Rescaling)     (None, 32, 32, 3)           0
conv2d (Conv2D)          (None, 32, 32, 64)          1792
max_pooling2d (MaxPooling2D) (None, 16, 16, 64)        0
conv2d_1 (Conv2D)         (None, 16, 16, 128)         73856
max_pooling2d_1 (MaxPooling2D) (None, 8, 8, 128)        0
```

✓ 0s completed at 1:54 AM

+ Code + Text

```
[16] epochs=35
numpy.random.seed(seed)
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=200)
258/258 [=====] - 4s 18ms/step - loss: 0.9850 - accuracy: 0.6528 - val_loss: 0.9766 - val_accuracy: 0.6665
Epoch 7/35
```

ICP5-16334245.ipynb - Colaboratory

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[18] val_acc = history.history['val_accuracy']
loss=history.history['loss']
val_loss=history.history['val_loss']

epochs_range = range(epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```

Training and Validation Accuracy

Training and Validation Loss

Training Loss Validation Loss

✓ 0s completed at 1:54 AM

```

  from keras.backend import clear_session
  clear_session()

[20] from keras.preprocessing.image import load_img
  from keras.preprocessing.image import img_to_array
  from keras.models import load_model
  from io import BytesIO
  from urllib.request import urlopen
  from PIL import Image
  import numpy as np

  def loadImage(URL):
    res = urlopen(URL).read()
    img = Image.open(BytesIO(res))
    fig = plt.figure(figsize=(10, 10))
    ax1 = fig.add_subplot(1, 2, 1)
    ax1.imshow(img)
    plt.axis('off')
    plt.title('Original')
    img = img.resize((32, 32))
    ax2 = fig.add_subplot(1, 2, 2)
    ax2.imshow(img)
    plt.axis('off')
    plt.title('Compressed to 32x32x3')
    fig.show()
    #convert to array
    mod = img_to_array(img)
    # reshape into a single sample with 3 channels
    mod = mod.reshape(1, 32, 32, 3)
    # prepare pixel data
    mod = mod.astype('float32')
    return mod

[22] def predict_class(img):
  cifar10_classes=['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
  prediction = model.predict(img)
  classes=np.argmax(prediction, axis=1)
  print(cifar10_classes[classes[0]])

  return mod

[23] deer = loadImage('https://media.istockphoto.com/photos/fallow-deer-picture-id1145651250?k=20&m=1145651250&s=612x612&w=0&h=F8zwduw-t7WCKX4HGTfe6nTomjoL7S2nrqrUs5JzvU=')

  Run cell (Mj/Ctrl+Enter)
  cell executed since last change (img):
  -['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
  executed by Mustavi Islam
  1:54 AM (4 minutes ago)
  del.predict(img)
  ax(prediction, axis=1)
  print(cifar10_classes[classes[0]])

  deer = loadImage('https://media.istockphoto.com/photos/fallow-deer-picture-id1145651250?k=20&m=1145651250&s=612x612&w=0&h=F8zwduw-t7WCKX4HGTfe6nTomjoL7S2nrqrUs5JzvU=')

  Compressed to 32x32x3


  [24] predict_class(deer)
  deer

  cat=loadImage('data:image/jpeg;base64,/9j/4AAQSkZJRgABAQAAAQABAD/2wCEAAoHCBYWFrgwFRRYGRg2GhoZGRkZHrkYHxoaGBwZHB0YGbgdIS41Hx4rIx0aJjomKy8xNTU1HCQ7QdszPy40NTEBDAwMEA8QHhISHj0rJCsNjQ0Nj0xMTQk
  Compressed to 32x32x3

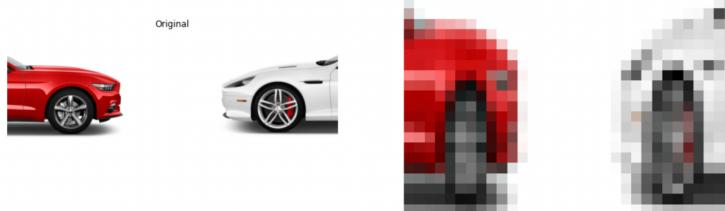


```

```
Q
<> Lets give it a hard image. The model was probably able to tell that the image had wheels and a hood, but it failed to predict 'automobile' over 'truck'
{x}
[27] car = loadImage('https://www.izmostock.com/wp-content/uploads/2018/04/izmostock_MainBanner_04.jpg')

```

Compressed to 32x32x3



```
✓ [28] predict_class(car)
↳ truck
```

```
<> [29] ship = loadImage('https://encrypted-tbn0.gstatic.com/images?q=tbn:ANd9GcS003YIEUhzgkXVMbpLZLFj0Vqg6QRH6nzGGw&usqp=CAU')
```

{x} Compressed to 32x32x3



```
✓ [30] predict_class(ship)
↳ cat
```

C

```
Q
<> [31] frog = loadImage('https://encrypted-tbn0.gstatic.com/images?q=tbn:ANd9GcQMG4yexr2e0fzLDq9C-VZGPUBL7B0wE219g&usqp=CAU')
```

Compressed to 32x32x3



```
✓ [32] predict_class(frog)
↳ frog
```