

# **University of Missouri-Kansas City**

## **School of Computing and Engineering**



**ICP:** ICP-05

**Course Name:** Big Data Analytics and Applications  
**Course ID:** COMP-SCI 5542

**Semester /Session :** Spring 2022

**Student ID : Student Name**

16334245 Mustavi Islam

16321217 Keenan Flynn

### **Description of the Problem :**

Recurrent neural networks and LSTMs help us solve problems with streamed data. This data can also be thought of as time series data. Each piece of information is fed into the RNN step by step. The goal of the RNN is to predict the next time step or the next series of time steps. RNNs and LSTMs do this by taking the previous outputs or hidden states as an input. In this way our network can begin to build up a type of memory so that it can predict certain types of sequences.

ICP 5 focuses on the layers and hyperparameter tuning of a recurrent neural network. We decided to make our model an LSTM as this type of model has higher memory capacity than a traditional LSTM. We also tuned some hyperparameters which can change the learning rate and computational cost of the model.

### **Description of Solution :**

1. The model provided in the source code for the ICP was adapted for a Project Gutenberg book
  - a. We used the book 'The Great Gatsby'. We sliced this initial book to get rid of some legal disclosures.
  - b. We added some preprocessing to the model. We used a Counter() object to aggregate the data into buckets by character and get counts for each character.
  - c. With these counts, we made some decisions to remove some characters and replace them with others. Essentially we removed outliers of the data. For instance, there was one instance of the character 'ç' which was replaced with a 'c'. We hope that this will have a small effect on the output.
  - d. We also assigned indexes based on these counts. There were 81 unique characters. The lowest frequency character 'Z' got assigned the index of 0 while the highest frequency character ' ' (space) got assigned the index of 80. By doing this we hope to introduce logic into the numeric indexes. The higher the number, the greater likelihood that the character has a high frequency.
2. Processing the text and changing sequence length.
  - a. We kept the text processing mostly the same as in the source code.
  - b. One change we made was to increase the sequence length. The sequence length is the amount of characters in a set which the model then has to predict from. We have a large dataset (270,000) characters, so increasing the sequence length will allow the model to have more information in which to make its decision.
  - c. We changed the sequence length to 200.

- d. Doing this gives the model 200 time steps in which to learn the characters.
3. Changing Batch size and creating an infinite dataset.
    - a. We learned in the last ICP that a low batch size and low epochs lead to a fast learning rate, but doing this fails to capture some of the variance in the data and thus increases the loss.
    - b. We decided to increase the batch size to 128. With a larger batch size, the model ‘sees’ more data before it starts learning. This allows the model to get closer to the global optima, meaning that the gradient is better.
    - c. With a larger batch size, we will need to increase the number of epochs as a higher batch size decreases the learning rate.
    - d. We also made our dataset ‘infinite’ through use of the repeat() helper function. This function makes it so we will never run out of data in the training process.
  4. Increasing epoch size
    - a. In this ICP we have increased the epoch size dynamically. Ideally we want the model to stop when it hits the most optimal epoch. By hard coding the epoch number we are unable to always stop at the best epoch.
    - b. One solution is to add callbacks. We have 2 callbacks, 1 tracks the weights of the previous epochs. The other stops the model when the loss ceases to decrease.
    - c. This way we can increase the epoch number to a size of our choosing, but the model will hopefully stop itself before then.
    - d. We used an epoch size of 50.
  5. Adding an LSTM layer
    - a. The LSTM layer replaced the GRU. The GRU has 2 gates: the reset and update gates, whereas the LSTM has 3 gates: input, output, and forget gates. The GRU is easy to stand up for simple problems, but the LSTM generally outperforms the GRU in language modeling problems because the LSTM has a longer ‘memory’.
    - b. The Keras model of sequential neural networks allowed us to easily replace the GRU layer with an LSTM layer.
    - c. To use the LSTM layer, we need an additional hyperparameter named the steps\_per\_epoch. Because we added the repeat() helper to the shuffle function, our dataset is theoretically endless. The steps\_per\_epoch parameter tells the model how many batches to learn on during each epoch. We found good performance by changing this number to 200.
  6. Adding dropout
    - a. During one of our test runs, the generate\_text() function generate an expert straight out of the training data. We want this function to generate

pseudo-real text and not repeat text that it has already seen. This is a case of overfitting.

- b. We can slow down overfitting by adding in a Dropout() layer so that some of the data becomes hidden from the LSTM.
- c. We used a dropout ratio of 0.1.

## 7. Changing Temperature

- a. During our testing, we found that sometimes the model started to copy text from the training set and output it as ‘new’ text. This is a case of overfitting. We found that changing the temperature in the generate\_text() function could help to prevent this from happening.
- b. Our final model has a temperature of 1.3.

## 8. Results

```
The Great Gatsby of New York had a problem. His friend Nick was in love with his wife. He decided to Walut the same people, behis beginning to look squarely at Christmas itself down into the room. But Wilson stood there a little bewing that was ...  
  
... One autumn night, five years before, they had been walking down the street when the leaver takes over.  
  
I don't know a sould had few vistaction that what look skid had worvied suddenly now the other car, and everyone, and you are, and you are, on it and his car."  
  
"Thy us sit didn't enter into the conversation, but long there. She couldn't be haughtif in his stand the young Quinns, divorced now," said Tom boys besiledlieve her handkerchief before a married "perspiration, that yacht never come well."  
  
"You said a bad driver, don't you worry object in July Gatsby's gotesque to as the familiar for Long Island, if I could get the entry. All I ask is that they should give me a start."  
  
"Ask Myrtle and the hummer."  
  
"I'm all right now that it was already behind.  
  
"Anywhing a little better off than he was, explained this ma
```

### Transcript of our output:

*The Great Gatsby of New York had a problem. His friend Nick was in love with his wife. He decided to Walut the same people, behis beginning to look squarely at Christmas itself down into the room. But Wilson stood there a little bewing that was ...*

*... One autumn night, five years before, they had been walking down the street when the leaver takes over.*

*I don't know a soul had few visitors that what look skid had worried suddenly now the other car, and everyone, and you are, and you are, on it and his car."*

*"They us sit didn't enter into the conversation, but long there. She couldn't be haughty in his stand the young Quinns, divorced now," said Tom boys besieledlieve her handkerchief before a married "perspiration, that yacht never come well."*

*"You said a bad driver, don't you worry object in July Gatsby's grotesque to as the familiar for Long Island, if I could get the entry.*

*All I ask is that they should give me a start."*

*"Ask Myrtle and the hummer."*

*"I'm all right now that it was already behind.*

*"Anything a little better off than he was, explained this ma*

### **Challenges:**

The biggest challenge this ICP faced was making informed decisions regarding the tuning of hyperparameters. The LSTM neural network takes a significant amount of time to train. We found that tuning these hyperparameters could drastically increase the running time of the model. Deadlines exist in this class and in industry settings, so doing research is extremely important before changing things about the model. This is where scientific method can greatly help expedite things. We need to first make observations about the model, form a hypothesis and ask a question, do research about the question, and finally test the hypothesis. In this way we are not wasting our time with poorly made hyperparameters.

### **Learning Outcomes :**

1. We learned how LSTMs work.
2. We learned how to process text data to help solve a time series question.
3. We learned about callback functions.

4. We learned how to test a neural network before training it so as to get an idea of the output.
5. We learned how loss was affected by batch and epoch size.

### Resources:

<https://towardsdatascience.com/generating-cooking-recipes-using-tensorflow-and-lstm-recurrent-neural-network-a7bf242acad3>

Used this link for general information and some code snippets

### Video Link:

<https://youtu.be/QMQ9mBdOfYE>

### Screenshots :

```

icp6 text gen.ipynb
PRO File Edit View Insert Runtime Tools Help All changes saved
Comment Share Settings RAM Disk Editing
+ Code + Text
Import our dependencies
import tensorflow as tf
tf.keras.backend.clear_session()
import numpy as np
import os
import time

Download the data
path_to_file = tf.keras.utils.get_file('gatsby.txt', 'https://www.gutenberg.org/cache/epub/64317/pg64317.txt')
Downloading data from https://www.gutenberg.org/cache/epub/64317/pg64317.txt
311296/306227 [=====] - 1s 2us/step
319488/306227 [=====] - 1s 2us/step

Read the data. In this ICP we are using the book 'The Great Gatsby' as our dataset.
Our data had some legal information attached to it, so we decided to slice that part off.
First, explore the text by looking at the length.

# Read, then decode for compatibility
text = open(path_to_file, 'rb').read().decode(encoding='utf-8')

# Slice text to get rid of all legal disclosures, we only want the raw text from the book
text = text[908:-18762]

# length of text is the number of characters in it
print('Length of text: {} characters'.format(len(text)))
Length of text: 277003 characters

# Take a look at the first 1000 characters in text
print(text[:1000])
The Great Gatsby
by
F. Scott Fitzgerald

```



```
We have changed the sequence length
We hope that by changing the sequence length to be longer, the LSTM will have more information to learn from.
Doing this increases the number of time steps in which the model has to make a prediction.

[✓] # The maximum length sentence we want for a single input in characters
seq_length = 100

examples_per_epoch = len(text)//seq_length+1
TensorSliceDataset: char_dataset
TensorSliceDataset with 277001 items is good for streaming data
char_dataset = tf.data.Dataset.from_tensor_slices(text_as_int)

print(type(char_dataset))

#take() displays first x members of the tensor DataSet. For each element, we can get the data by calling .numpy() method
# we can then enter i into our lookup array to get the resultant character
for i in char_dataset.take(10):
    print(idx2char[i.numpy()])
    print(i)

<class 'tensorflow.python.data.ops.dataset_ops.TensorSliceDataset'>
tf.Tensor(14, shape=(), dtype=int64)
tf.Tensor(14, shape=(), dtype=int64)
tf.Tensor(14, shape=(), dtype=int64)
tf.Tensor(88, shape=(), dtype=int64)
tf.Tensor(88, shape=(), dtype=int64)
tf.Tensor(88, shape=(), dtype=int64)
T
tf.Tensor(49, shape=(), dtype=int64)
h
tf.Tensor(71, shape=(), dtype=int64)
e
tf.Tensor(79, shape=(), dtype=int64)
tf.Tensor(88, shape=(), dtype=int64)
```

**icp6 text gen.ipynb**

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

Print the first examples input and target values:

Create training batches

We used tf.data to split the text into manageable sequences. But before feeding this data into the model, we need to shuffle the data and pack it into batches.

We added the .repeat() helper to the end of the shuffle method

This will allow the model to repeatable and so it will never end. This gives us a huge advantage as Deep Learning needs to have lots of data. This allows us to increase our Steps per epoch in the model.

```
[12] # Batch size
BATCH_SIZE = 64

# Buffer size to shuffle the dataset
# (TF data is designed to work with possibly infinite sequences,
# so it doesn't attempt to shuffle the entire sequence in memory. Instead,
# it maintains a buffer in which it shuffles elements).
BUFFER_SIZE = 10000

#Add the repeat() helper so that the dataset is endless
dataset = dataset.shuffle(BUFFER_SIZE).batch(BATCH_SIZE, drop_remainder=True).repeat()

dataset

<RepeatDataset element_spec=(TensorSpec(shape=(64, 100), dtype=tf.int64, name=None), TensorSpec(shape=(64, 100), dtype=tf.int64, name=None))>
```

**icp6 text gen.ipynb**

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

embedding\_dim dimensions;

tf.keras.layers.GRU: A type of RNN with size units=rnn\_units (You can also use a LSTM layer here.)

tf.keras.layers.Dense: The output layer, with vocab\_size outputs.

# Length of the vocabulary in chars
vocab\_size = len(vocab)

# The embedding dimension
embedding\_dim = 256

# Number of RNN units
rnn\_units = 1024

#KF: we could increase this

Adding a LSTM layer LSTM has a memory cell which can hold information in memory for a longer period of time. A set of gates is used to control when information enters the memory, when it's output, and when it's forgotten. We could have used the GRU but the problem with GRU is that it doesn't have the separate memory cell and they have fewer gate cells.

Adding a Dropout layer We have added a Dropout Layer of Dropout(0.1) which means that 10% of data will be dropped out which aims to decrease the possibility of overfitting, thus we can decrease the loss and make a better prediction.

```
[14] def build_model(vocab_size, embedding_dim, rnn_units, batch_size):
    model = tf.keras.Sequential([
        #The embedding layer helps to add semantic meaning to the input
        #With embedding we can do word math such as King - man + woman = Queen
        tf.keras.layers.Embedding(vocab_size, embedding_dim,
                                 batch_input_shape=[batch_size, None]),
        # The LSTM layer brings in the concept of cell state. These cell states are
        # how the network remembers what has previously been entered
        tf.keras.layers.LSTM(rnn_units,
                            return_sequences=True,
                            stateful=True,
                            recurrent_initializer='glorot_uniform'),
        # The dropout layer prevents overfitting
        tf.keras.layers.Dropout(0.1),
        # The dense layer is where the model decides what to output
        tf.keras.layers.Dense(vocab_size)
    ])
    return model
```

```

Here we build the model with the chosen params

[15]: model = build_model(
    vocab_size = len(vocab),
    embedding_dim=embedding_dim,
    rnn_units=rnn_units,
    batch_size=BATCH_SIZE)

Try the model

Now run the model to see that it behaves as expected.

First check the shape of the output:

[16]: for input_example_batch, target_example_batch in dataset.take(1):
    example_batch_predictions = model(input_example_batch)
    print(example_batch_predictions.shape, "# (batch_size, sequence_length, vocab_size)")

(64, 100, 81) # (batch_size, sequence_length, vocab_size)

In the above example, the sequence length of the input is XXXXXX but the model can be run on inputs of any length:

# We can summarize the model to get information about the layers
model.summary()

Model: "sequential"
+-----+-----+-----+
Layer (type)      Output Shape     Param #
+-----+-----+-----+
embedding (Embedding)    (64, None, 256)   20736
lstm (LSTM)        (64, None, 1024)   5246976
dropout (Dropout)   (64, None, 1024)   0
dense (Dense)      (64, None, 81)     83025
+-----+-----+-----+
Total params: 5,350,737
Trainable params: 5,350,737
Non-trainable params: 0

To get actual predictions from the model we need to sample from the output distribution, to get actual character indices. This distribution is defined by the logits over the character vocabulary.

Note: It is important to sample from this distribution as taking the argmax of the distribution can easily get the model stuck in a loop.

Try it for the first example in the batch:

[18]: # We have not actually run the model, so this information will be gibberish
# However it helps us to understand what is going on in the model
# We use random so that the model does not output the same thing for similar inputs
sampled_indices = tf.random.categorical(example_batch_predictions[:, 0], num_samples=1)
sampled_indices = tf.squeeze(sampled_indices, axis=-1).numpy()

#This gives us, at each timestep, a prediction of the next character index:
sampled_indices

array([27, 37, 68, 76, 9, 63, 54, 28, 71, 23, 32, 50, 64, 66, 19, 40, 48,
       57, 71, 67, 38, 66, 22, 29, 55, 80, 60, 7, 74, 29, 49, 56, 0, 58,
       41, 37, 22, 70, 54, 57, 43, 74, 35, 32, 31, 41, 71, 49, 12, 44, 1,
       12, 42, 79, 46, 58, 77, 44, 69, 70, 60, 80, 39, 50, 2, 77, 55, 16,
       46, 88, 58, 31, 14, 75, 22, 35, 23, 44, 80, 70, 72, 59, 8, 64, 19,
       52, 73, 21, 4, 19, 17, 6, 6, 78, 11, 67, 76, 28, 62, 41])

```

icp6 text gen.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share Settings RAM Disk Editing

Train the model

At this point the problem can be treated as a standard classification problem. Given the previous RNN state, and the input this time step, predict the class of the next character.

Attach an optimizer, and a loss function The standard `tf.keras.losses.sparse_categorical_crossentropy` loss function works in this case because it is applied across the last dimension of the predictions.

Because our model returns logits, we need to set the `from_logits` flag.

```

# Define the loss function, we use sparse because our data is not one hot encoded
def loss(labels, logits):
    return tf.keras.losses.sparse_categorical_crossentropy(labels, logits, from_logits=True)

# Before we run the model, lets see what the loss output will look like
example_batch_loss = loss(target_example_batch, example_batch_predictions)
print("Prediction shape: ", example_batch_predictions.shape, "# (batch_size, sequence_length, vocab_size)")
print("scalar_loss: ", example_batch_loss.numpy().mean())

Prediction shape: (64, 100, 81) # (batch_size, sequence_length, vocab_size)
scalar_loss: 4.39246

```

Configure the training procedure using the `tf.keras.Model.compile` method. We'll use `tf.keras.optimizers.Adam` with default arguments and the loss function.

```

[20]: # Compile the model with the adam optimizer
model.compile(optimizer='adam', loss=loss)

```

Configure checkpoints

Use a `tf.keras.callbacks.ModelCheckpoint` to ensure that checkpoints are saved during training:

```

Q Configure checkpoints
<> Use a tf.keras.callbacks.ModelCheckpoint to ensure that checkpoints are saved during training:
(x)
[21] # Directory where the checkpoints will be saved
checkpoint_dir = './training_checkpoints'
# Name of the checkpoint files
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt_{epoch}")

checkpoint_callback=tf.keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_prefix,
    save_weights_only=True)

```

We also define an early stopping callback. This callback will allow the model to stop if the loss does not keep improving.  
With this callback, we can use large number of epochs but the model will stop after the optimal loss is found.

```

[22] early_stopping_callback = tf.keras.callbacks.EarlyStopping(
    patience=5,
    monitor='loss',
    restore_best_weights=True,
    verbose=1
)

```

Execute the training  
To keep training time reasonable, use 10 epochs to train the model.

+ Code + Text

▼ Changing Epoch size from 10 to 30.

Increasing the epoch size aims to decrease the loss. The increased epoch size is to provide more iterations with our dataset.  
Increased epochs allow for more learning to be done.

```

[23] EPOCHS = 30
INITIAL_EPOCH = 1
STEPS_PER_EPOCH = 200

```

+ Code + Text

```

[24] Epoch 29/30
Epoch 29/30 [=====] - 12s 61ms/step - loss: 0.1812
Epoch 29/30 [=====] - 12s 61ms/step - loss: 0.1788
Epoch 30/30 [=====] - 12s 60ms/step - loss: 0.1761

```

Generate text  
Restore the latest checkpoint To keep this prediction step simple, use a batch size of 1.  
Because of the way the RNN state is passed from timestep to timestep, the model only accepts a fixed batch size once built.  
To run the model with a different batch\_size, we need to rebuild the model and restore the weights from the checkpoint.

```

tf.train.latest_checkpoint(checkpoint_dir)
model = build_model(vocab_size, embedding_dim, rnn_units, batch_size=1)

model.load_weights(tf.train.latest_checkpoint(checkpoint_dir))

model.build(tf.TensorShape([1, None]))
model.summary()

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(1, None, 256)	20736
lstm_1 (LSTM)	(1, None, 1024)	5246976
dropout_1 (Dropout)	(1, None, 1024)	0
dense_1 (Dense)	(1, None, 81)	83025

Total params: 5,356,737  
Trainable params: 5,356,737  
Non-trainable params: 0

NON-trainable params: 0

The prediction loop

The following code block generates the text:

It Starts by choosing a start string, initializing the RNN state and setting the number of characters to generate.

Get the prediction distribution of the next character using the start string and the RNN state.

Then, use a categorical distribution to calculate the index of the predicted character. Use this predicted character as our next input to the model.

The RNN state returned by the model is fed back into the model so that it now has more context, instead than only one character. After predicting the next character, the modified RNN states are again fed back into the model, which is how it learns as it gets more context from the previously predicted characters.

Looking at the generated text, you'll see the model knows when to capitalize, make paragraphs and imitates a Shakespeare-like writing vocabulary. With the small number of training epochs, it has not yet learned to form coherent sentences.

```

[25] def generate_text(model, start_string):
    # Evaluation step (generating text using the learned model)

    # Number of characters to generate
    num_generate = 1000

    # Converting our start string to numbers (vectorizing)
    input_eval = [char2idx[s] for s in start_string]
    input_eval = tf.expand_dims(input_eval, 0)

    # Empty string to store our results
    text_generated = []

    # Low temperatures results in more predictable text.
    # Higher temperatures results in more surprising text.
    # Experiment to find the best setting.
    temperature = 1.0

```

✓ 8s completed at 6:31 PM

+ Code + Text

```

[26] # Here batch size == 1
model.reset_states()
for i in range(num_generate):
    predictions = model(input_eval)
    # remove the batch dimension
    predictions = tf.squeeze(predictions, 0)

    # using a categorical distribution to predict the character returned by the model
    predictions = predictions / temperature
    predicted_id = tf.random.categorical(predictions, num_samples=1)[-1,0].numpy()

    # We pass the predicted character as the next input to the model
    # along with the previous hidden state
    input_eval = tf.expand_dims([predicted_id], 0)

    text_generated.append(idx2char[predicted_id])

return (start_string + ''.join(text_generated))

```

print(generate\_text(model, start\_string="The Great Gatsby of New York had a problem. His friend Nick was in love with his wife. He decided to "))

The Great Gatsby of New York had a problem. His friend Nick was in love with his wife. He decided to call to him. Miss Baker had been cold, but I meet around and more.

"I'm not like him," said Mr. Wolfshiem, shaking my hand earnestly, "and what do you think I did?"

"Go what?" demanded Tom.

"Two studies. One of them I call Michaelis had to ask that it was already behind him, somewhere back in that vast obscurity beyond the city, where there warn in are man can store up in his ghostly telaph, in a wasteback to play from his Proke to an contralction was distasteful to Gatsby.

"I don't think this afternoon."

"Very well. I'm too him that the thing that we remained to the book, reading each item aloud and then

✓ 8s completed at 6:31 PM